

UJWAL KANDI

☎ (737) 326-8598 📍 Austin, TX ✉ ujwal.kandi@outlook.com 🔗 linkedin.com/in/ujwalkandi 🌐 ujwalkandi.in/projects

Professional Summary

Data Operations and Observability Engineer with **4+ years** keeping production data pipelines and monitoring infrastructure running for Fortune 500 clients including Disney, Bank of America, and Marriott. I specialize in catching pipeline and data-quality issues before they reach customers, using structured incident management and **RCA** workflows that cut **Mean Time to Resolution (MTTR) by up to 75%** and held **99.8% SLA compliance** across **16+** enterprise environments. M.S. in Information Technology and Management, UT Austin.

Technical Skills

Languages: Python, Bash, SQL, YAML, XML

Monitoring: Grafana, Dynatrace, AWS CloudWatch, Azure Monitor, Prometheus, ELK Stack (Elasticsearch, Logstash, Kibana)

Pipelines & Data: Airflow, Matillion, Kafka, Spark, PySpark, AWS Glue, Databricks, PostgreSQL, MySQL, Redshift, SQLite

Cloud: AWS (EC2, S3, Lambda, Redshift, Glue, CloudWatch), Azure (ADLS, Log Analytics, APIM, Azure Monitor)

Infrastructure: Terraform, CloudFormation, Ansible, Docker, Kubernetes, Git, GitHub, Jenkins, GitLab CI/CD

Networking: TCP/IP, HTTP/S, DNS, REST APIs, SNMP, OSI Model

Tools: ServiceNow, PagerDuty, Jira, Confluence, Tableau, Power BI

Experience

Apple Inc. (Contract via Welo Data)

Jan 2025 – Present

Metadata Operations Engineer

Austin, TX

- Managed metadata ingestion and quality validation pipelines for Apple TV's global content catalog across **100+** international partner feeds, maintaining **sub-2s latency** and **99.9%** uptime through automated XML validation and rollback procedures
- Built **Grafana** dashboards and threshold-based alerts that flag feed and data-quality issues before publication, cutting feed failures by roughly **30%** and reducing detection time from hours to under **15 minutes**
- Standardized **Root Cause Analysis (RCA)** with reusable debugging runbooks for recurring batch-job failures, trimming resolution time by about **25%** and steadily driving down repeat incidents across the pipeline
- Designed failover and rollback procedures that keep content available through high-volume release windows, aligned with Apple's **data-governance** standards and service-level agreement (**SLA**) targets

Environment: Grafana, Apache Airflow, XML, AWS S3, Python, Bash, Jira, Confluence

Sports Excitement (Sports data analytics startup)

Aug 2024 – Jan 2025

Data Operations Engineer

Austin, TX

- Built and maintained **50+ daily Airflow ETL jobs** with DAG-level retries, SLA breach alerts, and automated failure recovery, holding around **99.8%** job reliability for the analytics team's downstream reporting
- Optimized Azure Data Lake partitioning and storage tiers, bringing query latency down from **~1.2s to ~550ms** and speeding up the downstream ML workloads that fed business dashboards
- Integrated Airflow alerting with **PagerDuty** so on-call engineers were paged within minutes of a failure, replacing manual log checks and cutting mean response time to under **5 minutes** against SLA targets
- Added pre-deployment data-quality and load tests that cut post-release pipeline failures by about **28%**

Environment: Apache Airflow, Azure ADLS, Azure Monitor, PagerDuty, Python, SQL, GitLab CI/CD, Power BI

Dover Fueling Solutions

Jun 2024 – Aug 2024

Site Reliability Engineer, Intern

Austin, TX

- Designed and built the internship's core deliverable: an offline-capable edge diagnostics system for fuel-station IoT devices that let field technicians detect and resolve faults on-site without escalation, cutting support tickets by about **35%** and keeping fault resolution under **10 minutes**
- Deployed a **3-pillar observability stack** combining Grafana dashboards, Azure APIM log analytics, and Kafka metrics exporters to enable centralized visibility across fuel station infrastructure; reduced issue detection from **15 min to 3 min** and **MTTR by 40%**
- Tuned alert thresholds and escalation rules to cut alert fatigue by roughly **45%** while preserving full coverage of high-severity (**P1/P2**) incidents

Environment: Grafana, Azure APIM, Apache Kafka, Python, SQLite, Docker, ServiceNow, Jira

Epsilon (Publicis Groupe)

Aug 2021 – Jul 2023

Production Support Engineer (promoted from Data Programmer I, Apr 2022)

Bengaluru, India

Production Support Engineer, Apr 2022 – Jul 2023

- Ran end-to-end **incident management** covering detection, triage, resolution, **RCA**, and prevention across **16** enterprise client environments including Disney, Bank of America, Unilever, and Marriott, holding **99.8% SLA compliance** and bringing **MTTR from 45 min down to 12 min**
- Built the observability stack on **AWS CloudWatch**, Kibana, and Dynatrace for metrics collection, log analysis, and distributed tracing, then tuned alert baselines and anomaly thresholds that caught roughly **60%** of issues before they hit customers
- Wrote **20+ runbooks**, SOPs, and postmortem documents and drove a blameless RCA culture, cutting repeat incidents by about **50%** and shortening on-call onboarding for new engineers
- Automated production pipeline failover with **AWS CloudFormation**, dropping **RTO from 2 hrs to 30 min** and holding **99.9% uptime** through critical processing windows

- Optimized **Matillion ETL** job scheduling across SLA windows, retry logic, and maintenance jobs, improving batch completion rate by **32%** and reducing operational overhead by **38%**
Data Programmer I, Aug 2021 – Apr 2022
- Built and maintained ETL pipelines for **16** global clients, reducing batch-failure frequency by **45%** and ensuring consistent data availability for mission-critical analytics platforms
- Deployed CloudWatch, Dynatrace, and Kibana observability stack and performed rapid data quality checks and emergency hotfixes, maintaining a **93%** client satisfaction rate across enterprise accounts
Environment: AWS CloudWatch, Dynatrace, Kibana, ELK Stack, Matillion, Apache Airflow, AWS S3, Redshift, CloudFormation, Python, SQL, Jenkins, Confluence

- Isthriwala (E-commerce operations startup)Jun 2020 – Jul 2021
- Data Management AnalystVijayawada, India
- Optimized SQL database operations for a platform serving **35,000+** customers and **100+ B2B clients** through index restructuring and query execution plan analysis, improving query response time by **~40%** and easing load on critical transactional systems
 - Built the **Tableau** KPI dashboards the management team used daily to track inventory and sales, lifting monthly performance targets by **15%**; identified growth opportunities and optimized customer retention, contributing to a **10%** increase in monthly subscriptions
Environment: MySQL, Tableau, Python, Excel

Projects

- A.I.D.E. – Observability and Automated RCA Engine — Prometheus, Grafana, Ollama, Python, LlamaStack
- Built a full-stack observability system that combines metrics analysis, log anomaly detection, and automated **RCA** workflows for API and infrastructure health, triaging common issues without manual investigation
 - Wired up **Prometheus** and **Grafana** for live dashboards and ran on-device Llama 3 inference through Ollama inside a **2GB** memory budget, hitting sub-second responses and cutting simulated **MTTR by 60%** (Meta ATX Llama Hackathon)
- FlipFinder Pro – Real-Time Data Pipeline and Anomaly Detection — Databricks, PySpark, Python
- Built a market-data ingestion and ML-scoring pipeline on **Databricks** for distributed processing and feature engineering, flagging statistical outliers across a high-volume live dataset
 - Hit **89%+** price-prediction accuracy and surfaced **\$20M+** in detected anomalies through a dashboard with profit, risk, and confidence scoring (Databricks Free Edition Hackathon)
- Azure Document Intelligence Pipeline — Azure AI Studio, Python (Capstone, Stronghold Investment Management)Jan–May 2024
- Built an automated document-processing pipeline using OCR and a Mixtral LLM to pull structured ownership data out of unstructured files and route verified records downstream
 - Cut document-processing costs by **20%** with end-to-end automation, after benchmarking accuracy, coherence, and latency across **4** Azure-hosted models to select the production one

Leadership and Achievements

- ElevenLabs Worldwide Hackathon – First Place and Best Product Built with n8nDec 2025
- Global Hackathon (1,000+ participants)Austin, TX
- Won **First Place** at a global hackathon of **1,000+** participants, building ARIA, a real-time voice-based retail intelligence agent using multi-agent LLM orchestration, ElevenLabs TTS, and n8n-powered automation for personalized recommendations and post-purchase workflows
- Seed Studio Ranger – AI/IoT Developer ProgramOct 2024 – Present
- AIoT Developer Educator ProgramAustin, TX
- Designed and delivered **10+** hands-on technical workshops on edge computing, IoT systems, and embedded AI to **100+** university students, building practical skills in AI-driven operational monitoring and edge diagnostics
- Dover Fueling Solutions Hackathon – First PlaceFeb 2024
- Annual Internal HackathonAustin, TX
- Won **First Place** designing an edge AI diagnostics system for fuel station operations; presented the solution to DFS executive leadership and industry stakeholders at the DFS User Conference 2024
- HACKSciCom 2022 – FinalistMay 2022
- Annual Hackathon, MIT Hacking Medicine co-hosted with RocheRemote
- Reached the finals building end-to-end **data pipelines** with a globally distributed team at one of healthcare’s premier hackathons, co-hosted by **MIT** and **Roche**

Education

- The University of Texas at Austin2023 – 2024
- M.S. in Information Technology and Management, McCombs School of BusinessAustin, TX
- KL University2016 – 2020
- B.Tech in Computer Engineering, Specialization in Data AnalyticsGuntur, India
- CGPA: 8.49/10.0 | Department Topper, Dean’s List