

UJWAL KANDI

☎ (737) 326-8598 📍 Austin, TX ✉ ujwal.kandi@outlook.com 🔗 linkedin.com/in/ujwalkandi 🌐 ujwalkandi.in/projects

Professional Summary

Data Operations and Observability Engineer with 4+ years keeping production data pipelines and monitoring infrastructure running for Fortune 500 clients including Disney, Bank of America, and Marriott. I specialize in catching pipeline and data-quality issues before they reach customers, using structured incident management and RCA workflows that cut **Mean Time to Resolution (MTTR) by up to 75%** and held **99.8% SLA compliance** across 16+ enterprise environments. M.S. in Information Technology and Management, UT Austin.

Technical Skills

Languages: Python, Bash, SQL, YAML, XML

Monitoring: Grafana, Dynatrace, AWS CloudWatch, Azure Monitor, Prometheus, ELK Stack (Elasticsearch, Logstash, Kibana)

Pipelines & Data: Airflow, Matillion, Kafka, Spark, PySpark, AWS Glue, Databricks, PostgreSQL, MySQL, Redshift, SQLite

Cloud: AWS (EC2, S3, Lambda, Redshift, Glue, CloudWatch), Azure (ADLS, Log Analytics, APIM, Azure Monitor)

Infrastructure: Terraform, CloudFormation, Ansible, Docker, Kubernetes, Git, GitHub, Jenkins, GitLab CI/CD

Networking: TCP/IP, HTTP/S, DNS, REST APIs, SNMP, OSI Model

Tools: ServiceNow, PagerDuty, Jira, Confluence, Tableau, Power BI

Experience

Apple Inc. (Contract via Welo Data)

Jan 2025 – Present

Metadata Operations Engineer

Austin, TX

- Managed metadata ingestion and quality validation pipelines for Apple TV's global content catalog across **100+ international partner feeds**, maintaining sub-2s latency and 99.9% uptime through automated XML validation and rollback procedures
- Built Grafana dashboards and threshold-based alerts that flag feed and data-quality issues before publication, **cutting feed failures ~30%** and reducing detection time from hours to under 15 minutes
- Standardized Root Cause Analysis (RCA) with reusable debugging runbooks for recurring batch-job failures, **trimming resolution time ~25%** and steadily driving down repeat incidents across the pipeline
- Designed failover and rollback procedures that keep content available through high-volume release windows, aligned with Apple's data-governance standards and service-level agreement (SLA) targets

Environment: Grafana, Apache Airflow, XML, AWS S3, Python, Bash, Jira, Confluence

Sports Excitement (Sports data analytics startup)

Oct 2024 – Jan 2025

Data Operations Engineer

Austin, TX

- Built and maintained 50+ daily Airflow ETL jobs with DAG-level retries, SLA breach alerts, and automated failure recovery, **holding ~99.8% job reliability** for the analytics team's downstream reporting
- Optimized Azure Data Lake partitioning and storage tiers, **cutting query latency from ~1.2s to ~550ms** and speeding up the downstream ML workloads that fed business dashboards
- Integrated Airflow alerting with PagerDuty so on-call engineers were paged within minutes of a failure, replacing manual log checks and **cutting mean response time to under 5 minutes** against SLA targets
- Added pre-deployment data-quality and load tests that **cut post-release pipeline failures ~28%**

Environment: Apache Airflow, Azure ADLS, Azure Monitor, PagerDuty, Python, SQL, GitLab CI/CD, Power BI

Dover Fueling Solutions

Jun 2024 – Aug 2024

Site Reliability Engineer, Intern

Austin, TX

- Designed and built the internship's core deliverable: an offline-capable edge diagnostics system for fuel-station IoT devices that let field technicians detect and resolve faults on-site without escalation, **cutting support tickets ~35%** and keeping fault resolution under 10 minutes
- Deployed a 3-pillar observability stack combining Grafana dashboards, Azure APIM log analytics, and Kafka metrics exporters to enable centralized visibility across fuel station infrastructure; **reduced issue detection from 15 min to 3 min** and MTTR by 40%
- Tuned alert thresholds and escalation rules to **cut alert fatigue ~45%** while preserving full coverage of high-severity (P1/P2) incidents

Environment: Grafana, Azure APIM, Apache Kafka, Python, SQLite, Docker, ServiceNow, Jira

Epsilon (Publicis Groupe)

Aug 2021 – Jul 2023

Production Support Engineer (promoted from Data Programmer I, Apr 2022)

Bengaluru, India

Production Support Engineer, Apr 2022 – Jul 2023

- Ran end-to-end incident management covering detection, triage, resolution, RCA, and prevention across 16 enterprise client environments including Disney, Bank of America, Unilever, and Marriott, **holding 99.8% SLA compliance** and **cutting MTTR from 45 min to 12 min**
- Built the observability stack on AWS CloudWatch, Kibana, and Dynatrace for metrics collection, log analysis, and distributed tracing, then tuned alert baselines and anomaly thresholds that **caught ~60% of issues before customer impact**
- Wrote 20+ runbooks, SOPs, and postmortem documents and drove a blameless RCA culture, **cutting repeat incidents ~50%** and shortening on-call onboarding for new engineers
- Automated production pipeline failover with AWS CloudFormation, **cutting RTO from 2 hrs to 30 min** and holding 99.9% uptime through critical processing windows

- Optimized Matillion ETL job scheduling across SLA windows, retry logic, and maintenance jobs, **improving batch completion 32%** and reducing operational overhead 38%
- Data Programmer I, Aug 2021 – Apr 2022*
- Built and maintained ETL pipelines for 16 global clients, **reducing batch-failure frequency 45%** and ensuring consistent data availability for mission-critical analytics platforms
 - Deployed CloudWatch, Dynatrace, and Kibana observability stack and performed rapid data quality checks and emergency hotfixes, **maintaining 93% client satisfaction** across enterprise accounts
- Environment:** AWS CloudWatch, Dynatrace, Kibana, ELK Stack, Matillion, Apache Airflow, AWS S3, Redshift, CloudFormation, Python, SQL, Jenkins, Confluence

Isthriwala (E-commerce operations startup)Jun 2020 – Jul 2021

Data Management AnalystVijayawada, India

- Optimized SQL database operations for a platform serving 35,000+ customers and 100+ B2B clients through index restructuring and query execution plan analysis, **improving query response time ~40%** and easing load on critical transactional systems
- Built the Tableau KPI dashboards the management team used daily to track inventory and sales, **lifting monthly performance targets 15%**; identified growth opportunities and optimized customer retention, **driving a 10% increase in monthly subscriptions**

Environment: MySQL, Tableau, Python, Excel

Projects

- A.I.D.E. – Observability and Automated RCA Engine** — *Prometheus, Grafana, Ollama, Python, LlamaStack*
- Built a full-stack observability system that combines metrics analysis, log anomaly detection, and automated RCA workflows for API and infrastructure health, triaging common issues without manual investigation
 - Wired up Prometheus and Grafana for live dashboards and ran on-device Llama 3 inference through Ollama inside a 2GB memory budget, hitting sub-second responses and **cutting simulated MTTR 60%** (Meta ATX Llama Hackathon)
- FlipFinder Pro – Real-Time Data Pipeline and Anomaly Detection** — *Databricks, PySpark, Python*
- Built a market-data ingestion and ML-scoring pipeline on Databricks for distributed processing and feature engineering, flagging statistical outliers across a high-volume live dataset
 - Hit 89%+ price-prediction accuracy** and **surfaced \$20M+ in detected anomalies** through a dashboard with profit, risk, and confidence scoring (Databricks Free Edition Hackathon)
- Azure Document Intelligence Pipeline** — *Azure AI Studio, Python (Capstone, Stronghold Investment Management)* **Jan–May 2024**
- Built an automated document-processing pipeline using OCR and a Mixtral LLM to pull structured ownership data out of unstructured files and route verified records downstream
 - Cut document-processing costs 20%** with end-to-end automation, after benchmarking accuracy, coherence, and latency across 4 Azure-hosted models to select the production one

Leadership and Achievements

- ElevenLabs Worldwide Hackathon – First Place and Best Product Built with n8n** **Dec 2025**
- Global Hackathon (1,000+ participants)* *Austin, TX*
- Won **First Place** at a global hackathon of 1,000+ participants, building ARIA, a real-time voice-based retail intelligence agent using multi-agent LLM orchestration, ElevenLabs TTS, and n8n-powered automation for personalized recommendations and post-purchase workflows
- Seed Studio Ranger – AI/IoT Developer Program** **Oct 2024 – Present**
- AIoT Developer Educator Program* *Austin, TX*
- Designed and delivered 10+ hands-on technical workshops on edge computing, IoT systems, and embedded AI to **100+** university students, building practical skills in AI-driven operational monitoring and edge diagnostics
- Dover Fueling Solutions Hackathon – First Place** **Feb 2024**
- Annual Internal Hackathon* *Austin, TX*
- Won **First Place** designing an edge AI diagnostics system for fuel station operations; presented the solution to DFS executive leadership and industry stakeholders at the DFS User Conference 2024
- HACKSciCom 2022 – Finalist** **May 2022**
- Annual Hackathon, MIT Hacking Medicine co-hosted with Roche* *Remote*
- Reached the finals building end-to-end data pipelines with a globally distributed team at one of healthcare’s premier hackathons, co-hosted by MIT and Roche

Education

- The University of Texas at Austin2023 – 2024
- M.S. in Information Technology and Management, McCombs School of BusinessAustin, TX
- KL University2016 – 2020
- B.Tech in Computer Engineering, Specialization in Data AnalyticsGuntur, India
- CGPA: 8.49/10.0 | Department Topper, Dean’s List